



**BUILD WITH CONFIDENCE**

# **A Deep Dive into Database Security with PostgreSQL**

**Roberto Mello**

**crunchydata**

# About me



- **Working with databases and PostgreSQL since 1999**
- **BS/MS in Computer Science from Utah State**
- **Principal Solutions Architect with Crunchy Data**
- **Outside of work: more tech stuff (raspberry pis, home automation, home servers, etc)**
- **Outside of tech: snowboarding, Austrian economics, scouting, travelling**

# Why Postgres

Established, Reliable  
and Secure

Feature Rich

Extensibility

No Central Owner

Hiring (Open Source)

## The Technical Details

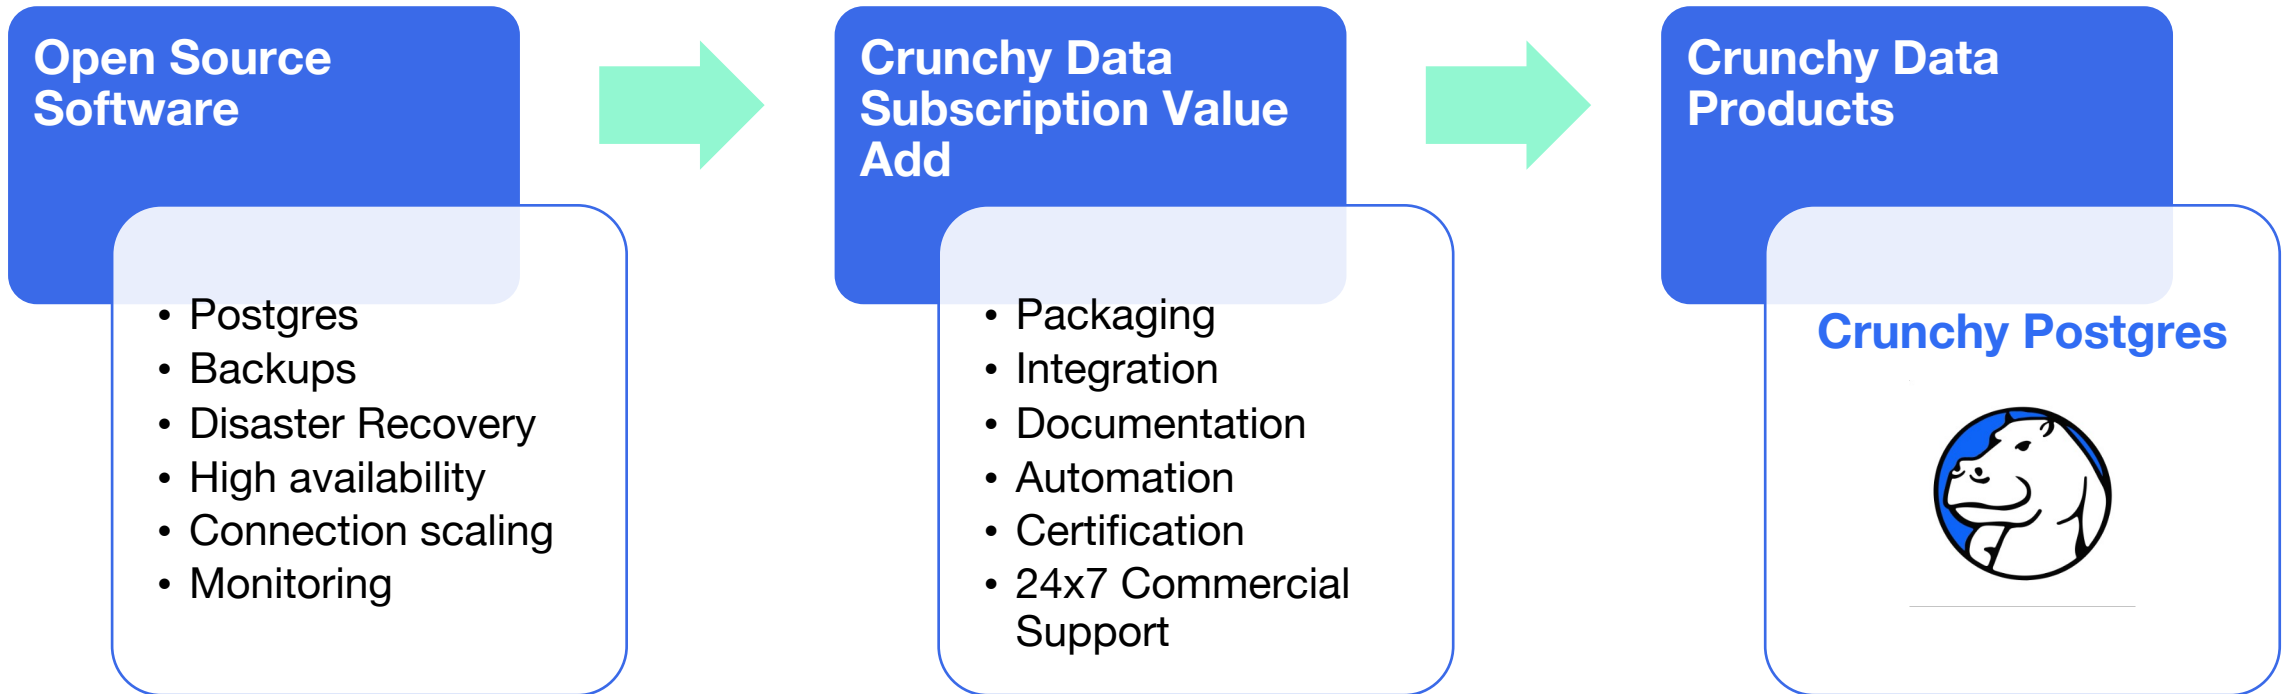
- Datatypes
- Transactional DDL
- Foreign Data Wrappers
- Concurrent Index Creation
- Conditional indexes
- JSON
- Common Table Expressions
- Fast column addition
- Listen/Notify
- Upsert
- Partitioning
- Per transaction sync replication
- Window function
- JSON/JSONB
- Continued innovation

# Crunchy Data



- **Leading Team in Postgres (core contributors)**
- **Certified Open Source Postgres Distribution**
- **Leader in Postgres Technology for Kubernetes**
- **Crunchy Bridge: Fully Managed Multi-Cloud Postgres**
- **Enterprise Focus and Go to Market Approach**

# Enterprise Ready Software Built on Open Source



# Crunchy Postgres is Production Ready Postgres



## Backups

Data is one of your most valuable assets, Crunchy Postgres ensures your data is safe and backed up so you can sleep easy.



## Disaster recovery

When the worst happens, we have you covered. Never worry about data corruption or loss with point-in-time recovery.



## High availability

With high availability you can trust your database is online.



## Connection scaling

With built in connection scaling you can easily scale to tens of thousands of connections for your database.



## Monitoring

Get the insights you need to know what is happening with your database with monitoring included.



## No lock-in

With Crunchy Data and open-source Postgres you are not locked-in to proprietary technology.

# Postgres Anywhere

## BARE METAL, VMS, CLOUD

### Crunchy Postgres

Crunchy Certified PostgreSQL is production ready Postgres.

#### INCLUDES:

- Backups
- Disaster recovery
- High availability
- Monitoring
- Automation
- Self managed

## KUBERNETES

### Crunchy Postgres for Kubernetes

Cloud Native Postgres on Kubernetes powered by Crunchy Postgres Operator.

#### INCLUDES:

- Simple provisioning
- Backups and DR included
- High availability
- Seamless upgrades
- Scale from 1 to thousands of databases
- Self managed

## FULLY MANAGED CLOUD

### Crunchy Bridge

The fully managed Postgres option on your choice of Cloud provider.

#### INCLUDES:

- AWS, Azure or GCP
- Continuous protection
- Backups
- Point in Time Recovery
- Full superuser access
- The developer experience you want

# Crunchy Bridge

## Cloud Postgres built for a superior developer experience

On Crunchy Bridge, Developer Experience (DX) means a toolset that gets you into the zone. Our platform has the tools necessary for developers to execute the optimum decision for performance and scale today while looking into the future.

*“The Crunchy team is awesome to work with, very responsive, very smart, and are clearly committed to getting the customer issues solved. We are working with Crunchy for the people and the technology.”* Seth Pollack, RivallQ Founder

The screenshot displays the Crunchy Bridge web interface for a cluster named 'mangy-tiger-693'. The interface is dark-themed with a blue header. The top navigation bar includes links for Overview, Networking, Backups, Logging, Container Apps, and Settings. Below the navigation bar, there are three summary cards: Storage (10 GB), CPU (1 Cores), and Memory (4 GB). The main section is titled 'Cluster Overview' and shows the database status as 'ready'. It includes a 'Credentials' section with a 'Role: Application' dropdown and a 'Choose format...' button. Below this, there is a table of cluster details: Provider (aws), Region (us-east-1), Plan (hobby-4), Version (Postgres 13), High Availability (No), and Disk Usage (Used 4.62 GB (29.48%), Total 15.67 GB, Available 10.23 GB (65.32%)). At the bottom, there is a 'Read Replicas' section with a table showing a single replica named 'mangy-tiger-693-replica' with details for Provider, Region, Plan, and Connection. A 'Copy URL' button is next to the connection details. The footer contains the Crunchy logo and links for Terms, Privacy, and Legal.

| REPLICA NAME            | PROVIDER | REGION         | PLAN    | CONNECTION               |
|-------------------------|----------|----------------|---------|--------------------------|
| mangy-tiger-693-replica | aws      | ap-southeast-2 | hobby-2 | <a href="#">Copy URL</a> |



# What is a Database?

- **An advanced SQL frontend to the file system**
- **A data access platform**
- **Daemon listening on (multiple) port(s)**
- **Files underneath**
  - **Local or remote storage?**
- **Transaction logs**
- **Backups**

# What is a Database?

```
# ss -tlnup | grep 5432
tcp    LISTEN 0          144          127.0.0.1:5432          0.0.0.0:*          users: (("postgres",pid=2260689,fd=5))
```


```
$ ls -la /var/lib/postgresql/16/main
```

```
total 88
drwx----- 19 postgres postgres 4096 Oct 11 18:32 .
drwxr-xr-x  3 postgres postgres 4096 Sep 18 11:52 ..
drwx-----  5 postgres postgres 4096 Sep 18 11:52 base
drwx-----  2 postgres postgres 4096 Oct 11 18:01 global
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_commit_ts
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_dynshmem
drwx-----  4 postgres postgres 4096 Oct 11 18:32 pg_logical
drwx-----  4 postgres postgres 4096 Sep 18 11:52 pg_multixact
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_notify
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_replslot
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_serial
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_snapshots
drwx-----  2 postgres postgres 4096 Oct 11 18:32 pg_stat
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_stat_tmp
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_subtrans
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_tblspc
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_twophase
-rw-----  1 postgres postgres    3 Sep 18 11:52 PG_VERSION
drwx-----  3 postgres postgres 4096 Sep 18 11:52 pg_wal
drwx-----  2 postgres postgres 4096 Sep 18 11:52 pg_xact
-rw-----  1 postgres postgres   88 Sep 18 11:52 postgresql.auto.conf
-rw-----  1 postgres postgres  130 Oct 11 18:01 postmaster.opts
```

# Database Security

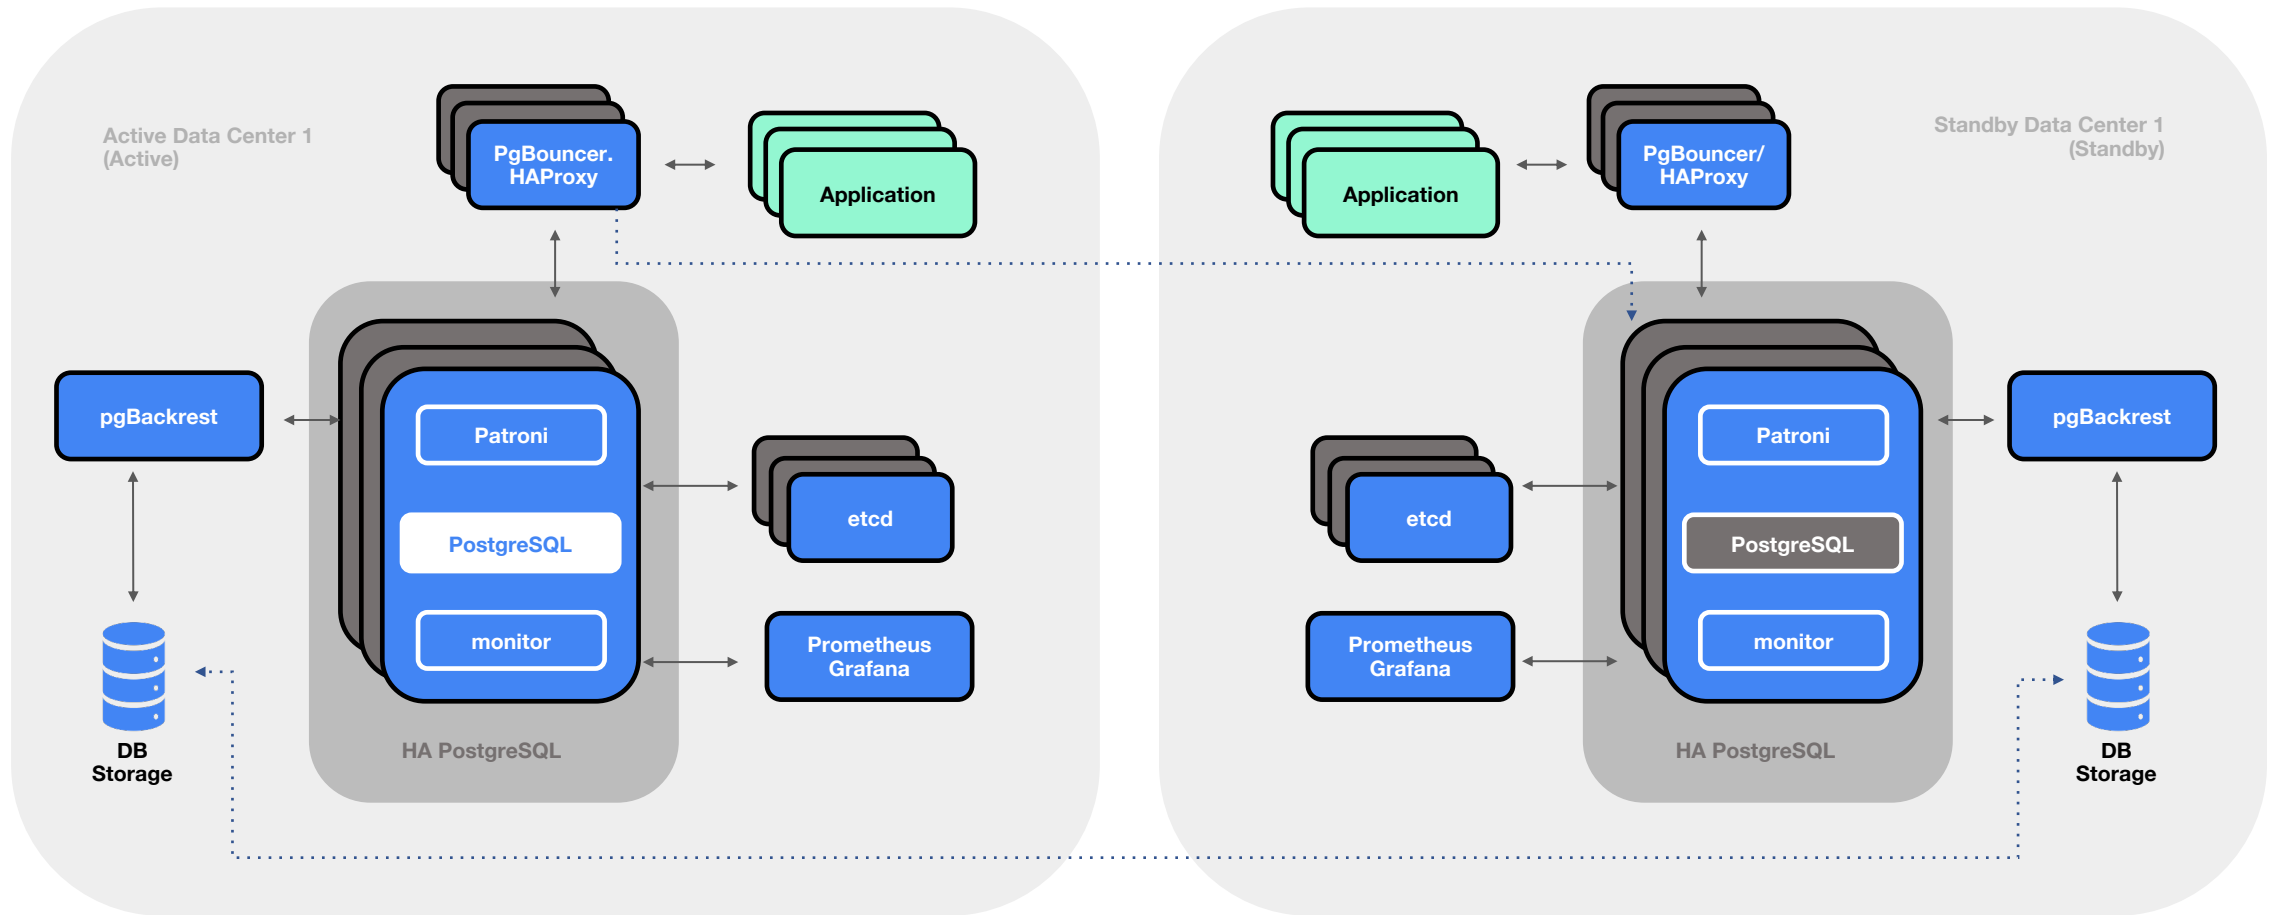
- **Who can access the database?**
- **Who can authenticate?**
- **What data can they access?**
- **What data can they change?**
- **Who did what when? (auditing)**

# Attack Vectors – simple view




- Who can access the database?
- Who can authenticate?
- What data can they access?
- What data can they change?
- Who did what when? (auditing)
- How can we recover?

# Attack Vectors – Multi Data Center with HA/DR



# External Factors



- **Pre-authentication**
- **Physical security**
- **Network security**
- **Sideway attacks**
  - Backups
  - Applications
- **Virtualization attacks**

# Who can Authenticate to the DB?

- **Principle of Least Privilege**
  - IP ranges
  - Users/Roles
- **Users and Roles in PostgreSQL**
  - A user is a role with login privilege
  - Users and Groups are the same
- **Use different users**
  - Very low-level user with minimal privileges
  - Strong user to do more, like create/drop tables
  - Superuser should be avoided (just like root)
- **Always use SSL**

# Who can Authenticate to the DB?

- **Host-based authentication (pg\_hba.conf)**
- **Which users can connect...**
- **...to which databases**
- **...connecting from which IP ranges (or locally)**
- **...with (non)encrypted connections**
- **...using which authentication methods**



# Who can Authenticate to the DB?

- **pg\_hba.conf**
- **Reloaded via SIGHUP or “SELECT pg\_reload\_conf()”**

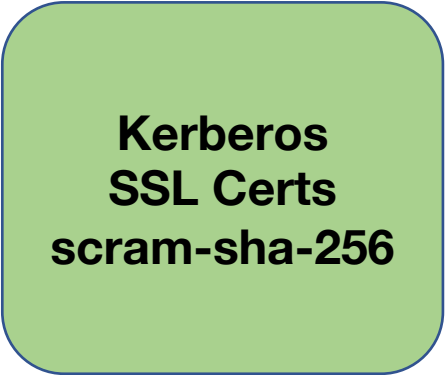
```
# TYPE      DATABASE          USER                ADDRESS              METHOD
# Database administrative login by Unix domain socket
local      all                postgres            peer
# "local" is for Unix domain socket connections only
local      all                all                 peer

# IPv4 / IPv6 local connections
host       app1                user_foo            127.0.0.1/32         scram-sha-256
host       app2                +group_bar          ::1/128              scram-sha-256

# SSL connections with regex matching databases and users (PG 16)
hostssl    "/^db\d{1,4}$" "/^user\d{1,4}$" 10.10.2.0/24         scram-sha-256
```

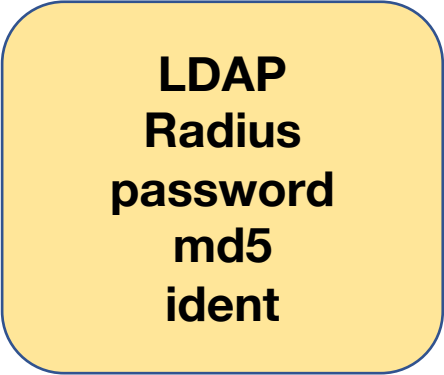
# Who can Authenticate to the DB?

- **Many companies deploy some sort of directory**
  - **Active Directory**
  - **FreeIPA**
- **Common mistake is to use LDAP authentication instead of Kerberos**



Kerberos  
SSL Certs  
scram-sha-256

Good



LDAP  
Radius  
password  
md5  
ident

Bad



trust  
pam

Ugly

# Who can Authenticate to the DB?

- **scram-sha-256 is a \*big\* improvement over md5**
  - does not reveal the user's cleartext password to the server
  - is designed to prevent replay attacks
- **SCRAM's Channel Binding is a feature that allows for mutual authentication (client/server) to prevent MITM attacks**
- **For extra security, require SCRAM channel binding**
- `channel_binding=require` **in the connection string**
- `PGCHANNELBINDING=require` **environment variable**

<https://www.citusdata.com/blog/2020/07/28/securely-authenticate-with-scram-in-postgres-13/>

# What can they see or do?

```
CREATE ROLE name [ [ WITH ] option [ ... ] ]
```

where *option* can be:

|  |  |
|--|--|
| SUPERUSER   NOSUPERUSER                                    | # Avoid superusers                         |
| CREATEDB   NOCREATEDB                                      | # Can it create databases?                 |
| CREATEROLE   NOCREATEROLE                                  | # Can it create other roles?               |
| INHERIT   NOINHERIT  | # PG 16: GRANT WITH INHERIT true           |
| LOGIN   NOLOGIN  | # Can it login? (user)                     |
| REPLICATION   NOREPLICATION                                | # Role will have replication privs         |
| BYPASSRLS   NOBYPASSRLS                                    | # Bypass or not Row Level Security         |
| <b>CONNECTION LIMIT</b> <i>conlimit</i>                    | <b># Default is no limit (-1)</b>          |
| [ ENCRYPTED ] PASSWORD ' <i>password</i> '   PASSWORD NULL |  |
| VALID UNTIL ' <i>timestamp</i> '                           | # Cut off access after this timestamp      |
| IN ROLE <i>role_name</i> [, ...]                           | # Added as member of existing roles        |
| IN GROUP <i>role_name</i> [, ...]                          | # Same as "IN ROLE"                        |
| ROLE <i>role_name</i> [, ...]                              | # Add existing roles into this new one     |
| ADMIN <i>role_name</i> [, ...]                             | # Add existing roles into this, being able |
|  | # add new roles into it (WITH ADMIN)       |

# What can they see or do?

## Several built-in roles

```
\du pg*
```

| List of roles               |              |
|-----------------------------|--------------|
| Role name                   | Attributes   |
| -----                       | -----        |
| pg_checkpoint               | Cannot login |
| pg_create_subscription      | Cannot login |
| pg_database_owner           | Cannot login |
| pg_execute_server_program   | Cannot login |
| pg_monitor                  | Cannot login |
| pg_read_all_data            | Cannot login |
| pg_read_all_settings        | Cannot login |
| pg_read_all_stats           | Cannot login |
| pg_read_server_files        | Cannot login |
| pg_signal_backend           | Cannot login |
| pg_stat_scan_tables         | Cannot login |
| pg_use_reserved_connections | Cannot login |
| pg_write_all_data           | Cannot login |
| pg_write_server_files       | Cannot login |

# What can they see or do?

- **Create ROLES with appropriate permissions to database objects**
- **Databases, schemas, tables, views, functions, sequences**
- **GRANT those roles to specific users**
- **GRANTs can be column based**
- **Remember Principle of Least Privilege**

# What can they see or do?

```
GRANT SELECT ON TABLE table_name TO ro_role;  
GRANT INSERT, UPDATE, DELETE ON TABLE table_name TO rw_role;  
GRANT EXECUTE ON FUNCTION function_name(argument_types) TO role;  
GRANT SELECT ON ALL TABLES IN SCHEMA schema_name TO role;  
GRANT USAGE ON SCHEMA schema_name TO role;  
GRANT CONNECT ON DATABASE database_name TO role;  
GRANT TEMP ON DATABASE database_name TO role;  
GRANT USAGE, SELECT ON SEQUENCE sequence_name TO role;  
GRANT USAGE, SELECT ON ALL SEQUENCES IN SCHEMA schema_name TO role;  
GRANT role TO another_role;
```

## GRANT can be column based

```
GRANT SELECT (column_name), UPDATE (column_name) ON TABLE  
table_name TO role;
```

# Security Definer function/stored proc

- Resist the temptation to make a role a superuser
- If you need a complicated set of permissions, create a stored procedure with **SECURITY DEFINER** privilege
- Be careful to avoid allowing a security hole (see docs)



# Security Definer function/stored proc

-- As user with elevated privileges



```
CREATE OR REPLACE FUNCTION insert_sensitive_data(data text) RETURNS void AS
$$
BEGIN
    -- Check if the current user has the required privileges
    IF current_user = 'trusted_user' THEN
        INSERT INTO sensitive_data (data_value) VALUES (data);
    ELSE
        RAISE EXCEPTION 'Permission denied.';
    END IF;
END;
$$ LANGUAGE plpgsql SECURITY DEFINER;
```

-- As regular user: Grant execute privilege to a trusted user

```
GRANT EXECUTE ON FUNCTION insert_sensitive_data(text) TO trusted_user;
```

# What can they see or do?

Understand the chain of objects and their permissions



# Row Level Security

- **In Postgres since 2014**
  - **MySQL doesn't have it. Views only**
- **Fine-grained Access Control**
- **Restrict, on a per-user basis**
  - **Which rows are visible to normal queries (per user)**
  - **What can be inserted, updated, deleted**
- **Policies control visibility**
- **Used alongside GRANT/REVOKE**

# Row Level Security

```
SET ROLE postgres;
```

```
CREATE POLICY security_clearance ON hero
  USING (pg_has_role(current_user, clearance, 'MEMBER'));
```

```
ALTER TABLE hero ENABLE ROW LEVEL SECURITY;-- Everyone but superusers+owner
ALTER TABLE hero FORCE ROW LEVEL SECURITY; -- Everyone but superusers
```

| Table "public.hero" |      |           |          |         |  |
|---------------------|------|-----------|----------|---------|--|
| Column              | Type | Collation | Nullable | Default |  |
| name                | text |           |          |         |  |
| secret_identity     | text |           |          |         |  |
| powers              | text |           |          |         |  |
| weaknesses          | text |           |          |         |  |
| clearance           | text |           |          |         |  |

Policies (forced row security enabled):

```
POLICY "security_clearance"
  USING (pg_has_role(CURRENT_USER,
    (clearance)::name, 'MEMBER'::text))
```

# Mandatory Access Control

- **sepgsql**
  - SELinux bindings
- **RBAC type enforcement covers most DB objects**
- **Can be combined with custom SELinux policy for powerful control**

# What about data on disk?

- **Transparent Data Encryption is still a WIP in Postgres, available in closed products**
- **Pgcrypto – extension to encrypt data in the database**
  - PGP, OpenSSL; hashing and encryption
  - Not very practical
  - Do not save the key next to the data itself!
- **PostgreSQL Anonymizer**
  - Extension to mask or anonymize PII in the database
- **Outside of PostgreSQL**
  - Operating System
  - Storage device / cloud storage

# What about the superuser?

- **Bypass all DAC (Discretionary Access Control)**
- **Bypass all Row Level Security**
- **Can load any library**
- **Can run COPY...PROGRAM (execute arbitrary shell cmds)**
- **Can run ALTER SYSTEM (change conf settings with SQL)**
- **Create/execute any function**
- **Others**

# set\_user extension

- **GRANT EXECUTE on set\_user() and/or set\_user\_u() to otherwise unprivileged users**
- **Can switch the effective user when needed to perform specific actions**
- **Optional enhanced logging ensures an audit trail**
- **Once one or more unprivileged users able to run set\_user\_u(), ALTER superuser to NOLOGIN**
- **Multiplex unprivileged users, e.g. with connection pools**



# Who did what, and when?

- **Database logs**
  - Ship them somewhere else (or outsource)
  - Have a sensible logging policy (don't log connections)
- **PostgreSQL parameters**
  - `log_connections`
  - `log_statement (ddl, mod, all)` – beware in production
- **Log Rotation/truncation/removal**
- **PGAudit extension**
  - Detailed session and/or object logging
  - Who did exactly what, and when
  - Produces audit logs required for compliance

# What Else Can Go Wrong?

- **Keep server updated to the latest minor release**
  - Only bug fixes. Don't hesitate to upgrade.
- **Extensions have their own upgrade/patching cycle**
- **Ecosystem software (pgbackrest, patroni, pgbouncer)**
  - needs to match versions of Postgres they work with
  - have their own release cycle
- **Do not delete WAL files**

# Don't Forget Backups

- **Database dumps are NOT backups**
- **Storage snapshots are NOT (reliable) backups**
- **Everybody needs Point In Time Recovery (PITR) – why?**
- **DB “Backups” are comprised of separate operations**
  - **Restore**
  - **Recovery – replaying of transaction logs up to a consistent state**
- **If you're not exercising \*documented\* restores, you can't assume your backups are valid**

# Don't Forget Backups

- Use specialized software like pgBackRest
  - Backup encryption
  - multi-repository, delta restores, file bundling, block incremental
- Last words heard in the operations room:

**“But I thought \*you\* did the backups!”**



Casually  
executing a  
delete  
statement

You forgot the  
where clause

# Should I run PostgreSQL in a container?

**You already are running PostgreSQL in a container**



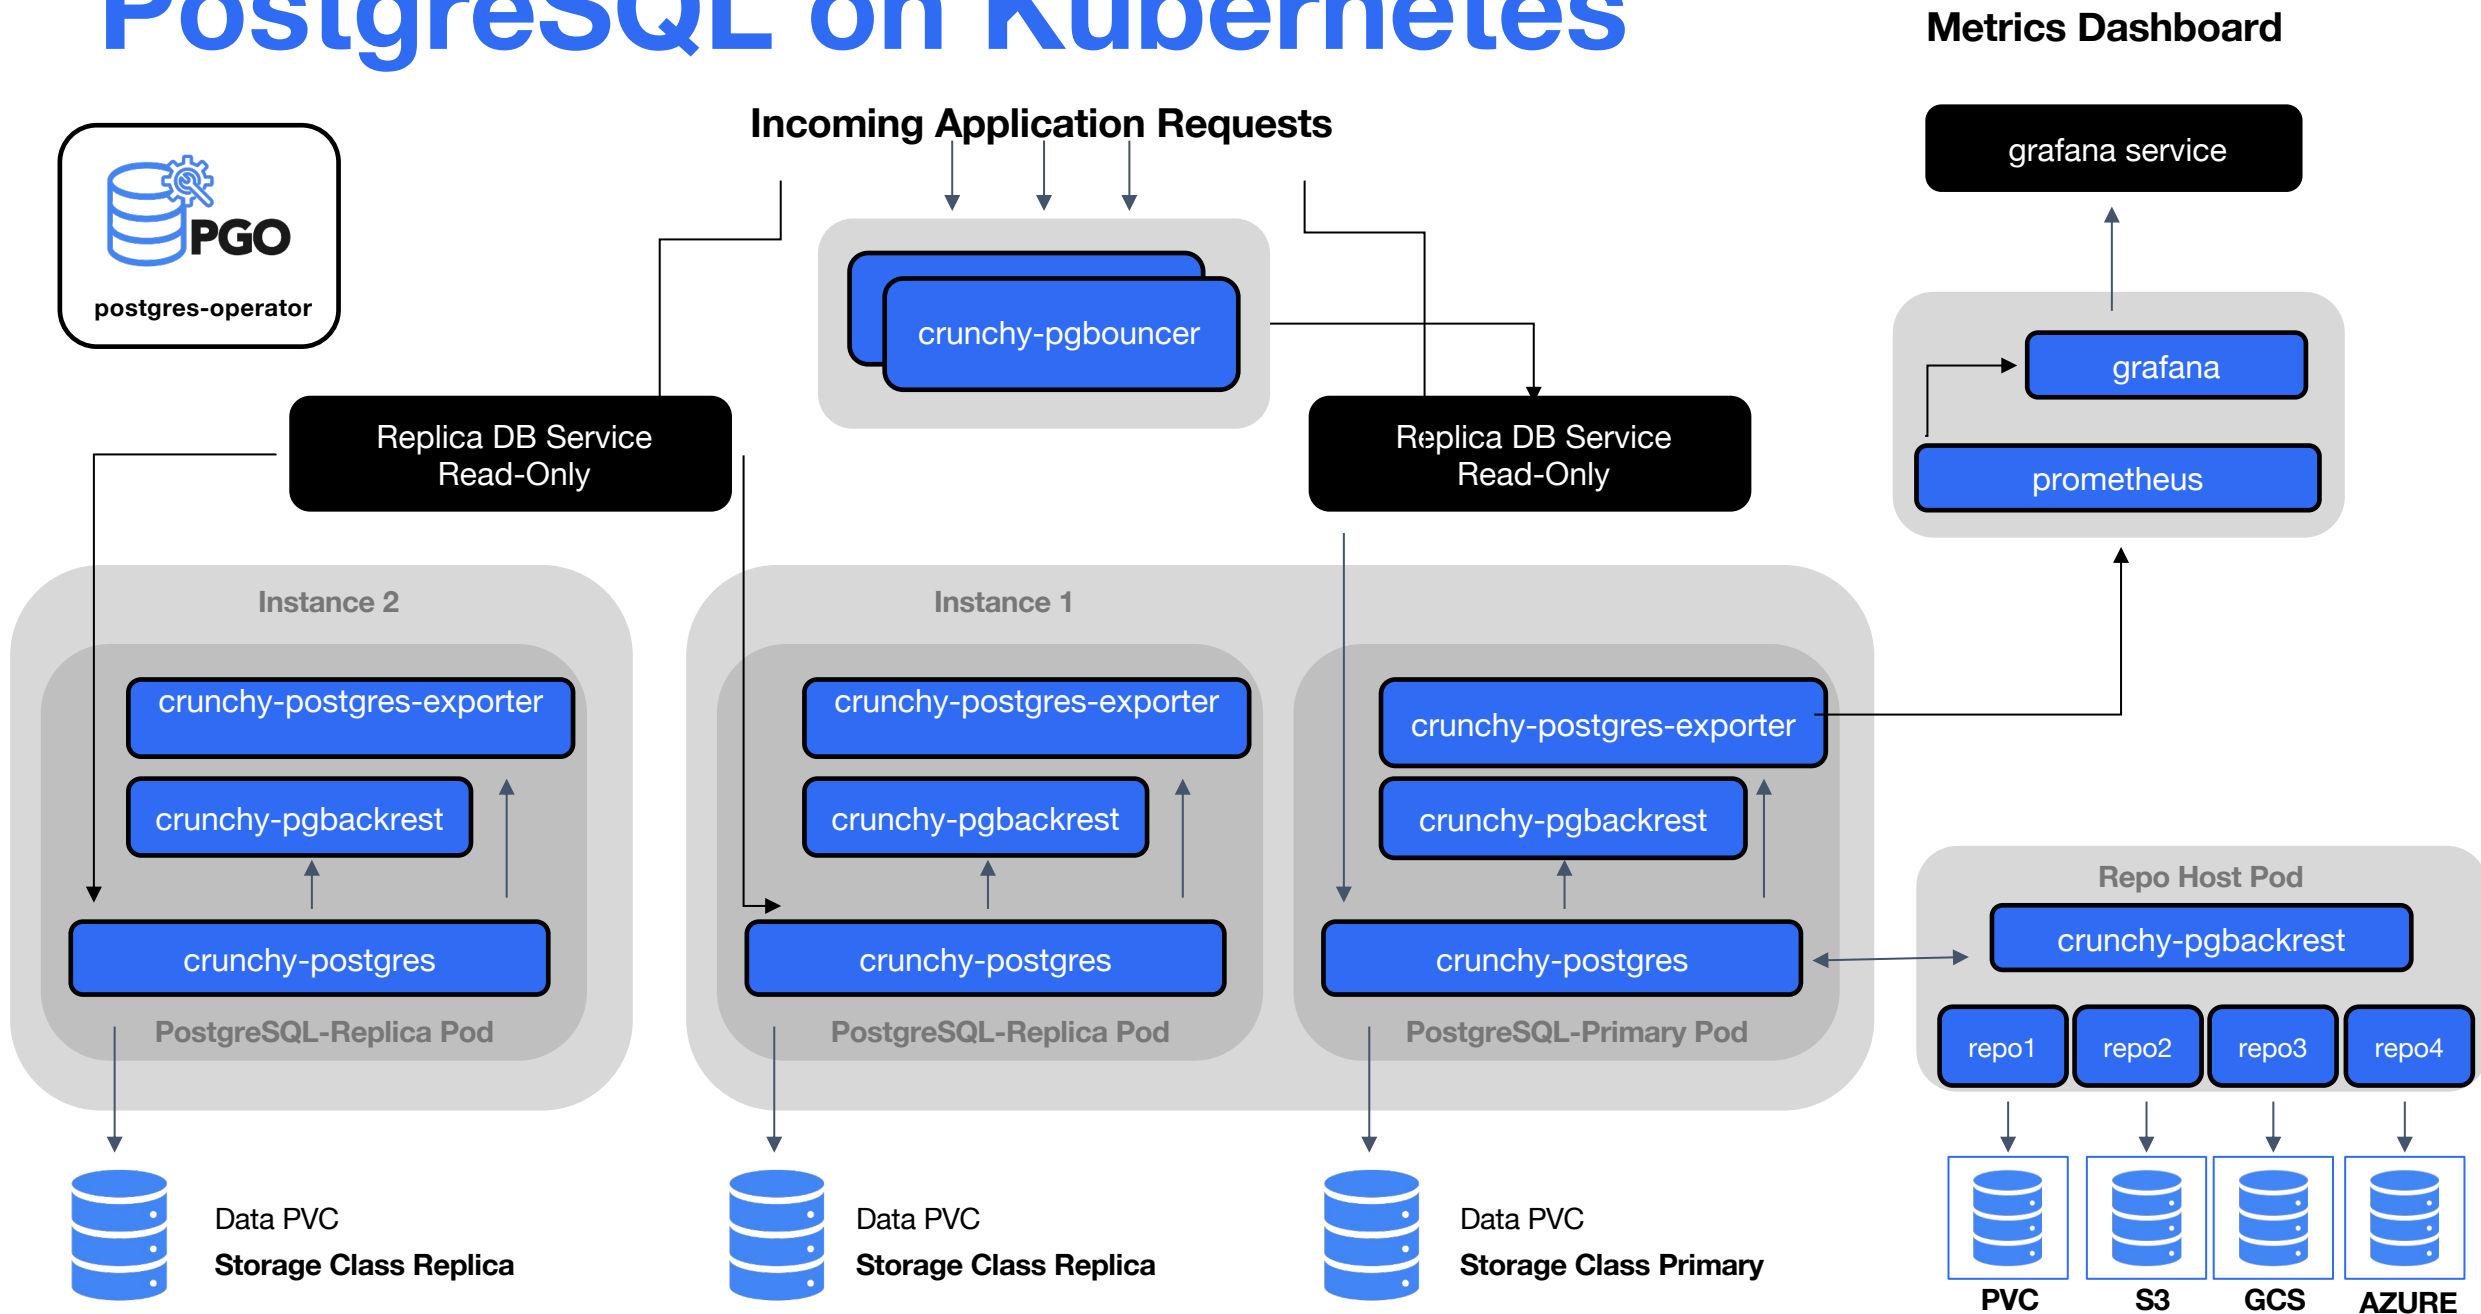
```
# ps -fae | grep postgres
```

```
postgres 248985      1  0 Feb22 ?           00:00:02 /usr/pgsql-13/bin/postgres -D  
/app/pgdata/hippo.13  
...
```

```
# 11 /proc/248985/ns
```

```
total 0  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 12:59 cgroup -> 'cgroup:[4026531835]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 12:59 ipc -> 'ipc:[4026531839]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 12:59 mnt -> 'mnt:[4026531840]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 12:59 net -> 'net:[4026531992]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 12:59 pid -> 'pid:[4026531836]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 13:11 pid_for_children -> 'pid:[4026531836]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 13:11 time -> 'time:[4026531834]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 13:11 time_for_children -> 'time:[4026531834]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 12:59 user -> 'user:[4026531837]'  
lrwxrwxrwx 1 postgres postgres 0 Feb 24 12:59 uts -> 'uts:[4026531838]'
```

# PostgreSQL on Kubernetes



# It is a container, what could go wrong?

## Managing compute resources

**Overprovisioning**

**Control Plane**

**No Monitoring / Capacity Planning**

**Proper Storage**

## Container resource limits

**Minimal CPU\*:**

$\text{<number of active concurrent sessions> / 2.5}$

**Minimal Memory\*:**

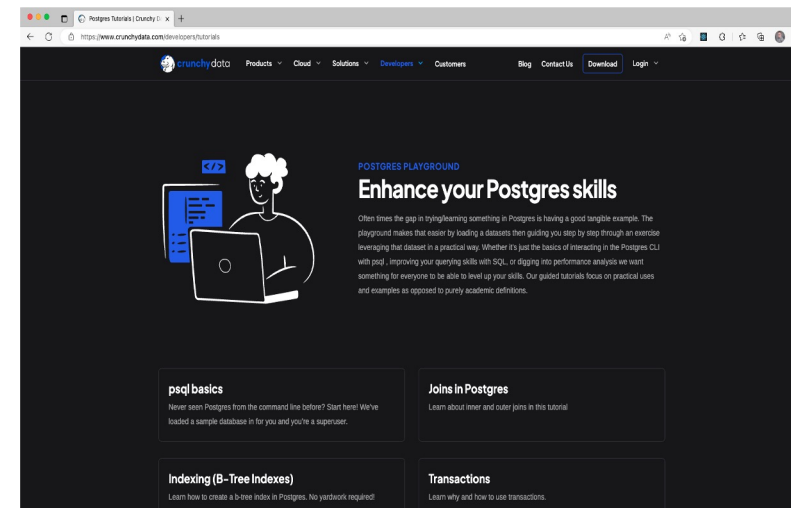
$(\text{work\_mem} * \text{<avg active concurrent sessions>} * 2) +$   
 $(\text{maintenance\_work\_mem} * \text{autovacuum\_max\_workers})$   
 $+ \text{shared\_buffers} + (\text{max\_connections} * 20\text{MB}) < 70\%$



## Immature Kubernetes Process/Procedures

# Acknowledgements and Resources

- **Greg Sabino Mullane**
- **Joe Conway**
- **Crunchy Data**
- **All of the amazing PostgreSQL community**
- **Run and Learn PostgreSQL in your browser**  
<https://www.crunchydata.com/developers/tutorials>
- **Learn PostgreSQL**  
<https://www.crunchydata.com/developers/>
- **PGPedia:** <https://pgpedia.info/>







**BUILD WITH CONFIDENCE**

**Thank you**